

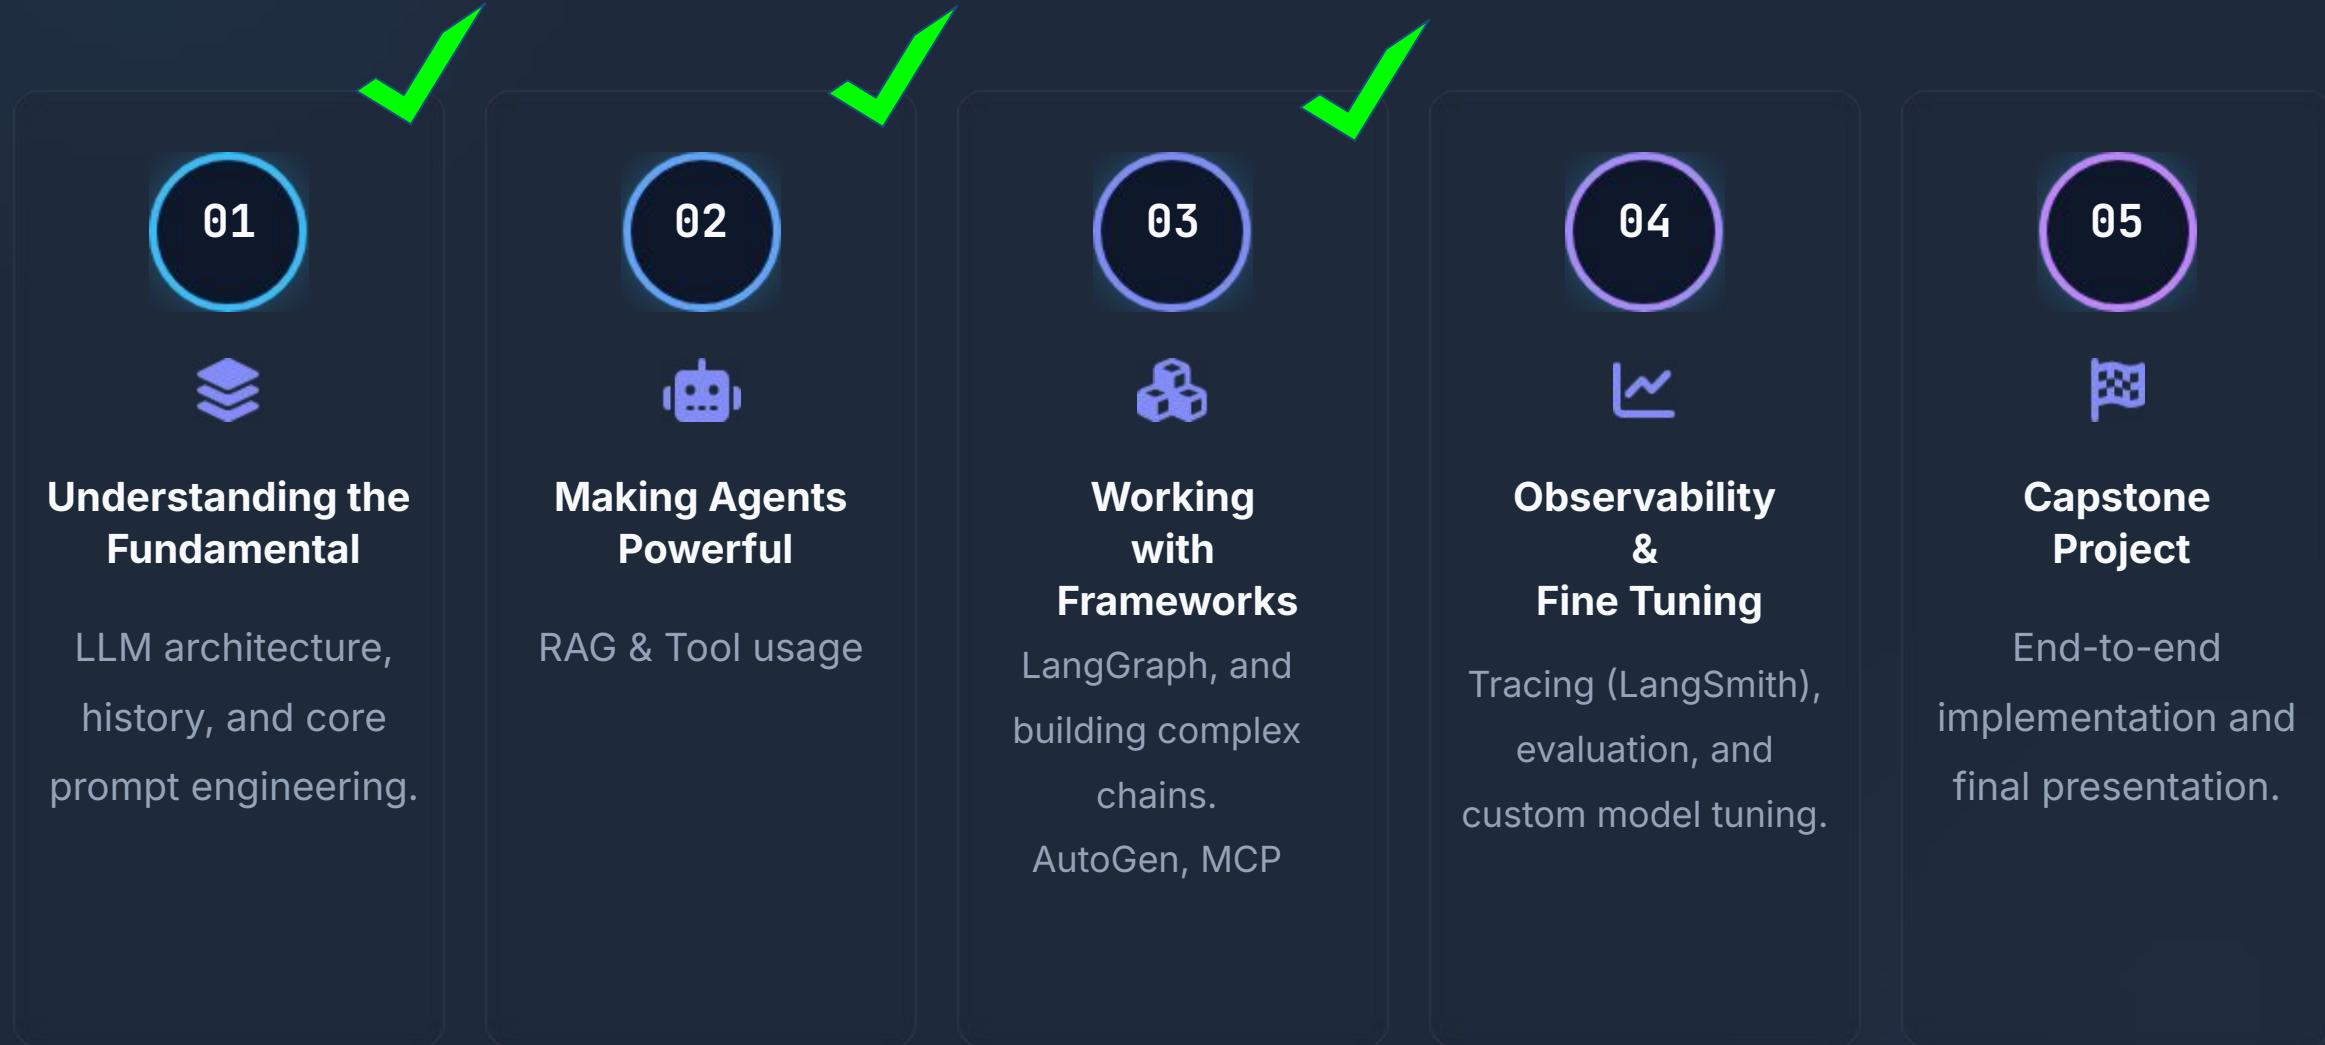
**DAY 4: LLMOPS & MODEL CUSTOMIZATION**

# Adding Observability and Fine-Tuning

Ensuring Reliability, Traceability, and Domain Expertise

# Course Roadmap: From Foundations to Capstone

---



# **Day 5 Mini Project Details**

- **Pick a project**
- **Create team of 3-4 members**
- **Follow the guidelines given in document**
- **Each team will need to present the solution**

# Day 4 Roadmap

---

## Key Focus Areas

Today is dedicated to the infrastructure required to move an LLM prototype into a robust enterprise application.



**Agentic RAG Pipeline** Implementing Planner → Retriever → Answerer flow ( RAG Github Repositories Overview)



**Practical Fine-Tuning** Using adapters (PEFT) for domain specificity and comparing outputs



**Evaluation Strategies:** How to evaluate LLM Applications



**LLMOPs Foundations (Tracing)** Logging prompts, tracking latency/cost, and tracing workflows, PII redaction layers

# Langchain Tool usage with Autogen

# Task 1

- Multi Tool calling Agent using API

# Task 2

- Enhancement in SQL Agent

# Agentic RAG

## **Agentic RAG essentially combines**

- **RAG (Retrieval-Augmented Generation):** The technique of retrieving relevant information from external knowledge sources to augment LLM responses
- **Agentic AI patterns:** Giving the AI autonomous decision-making capabilities to plan, reason, and take actions

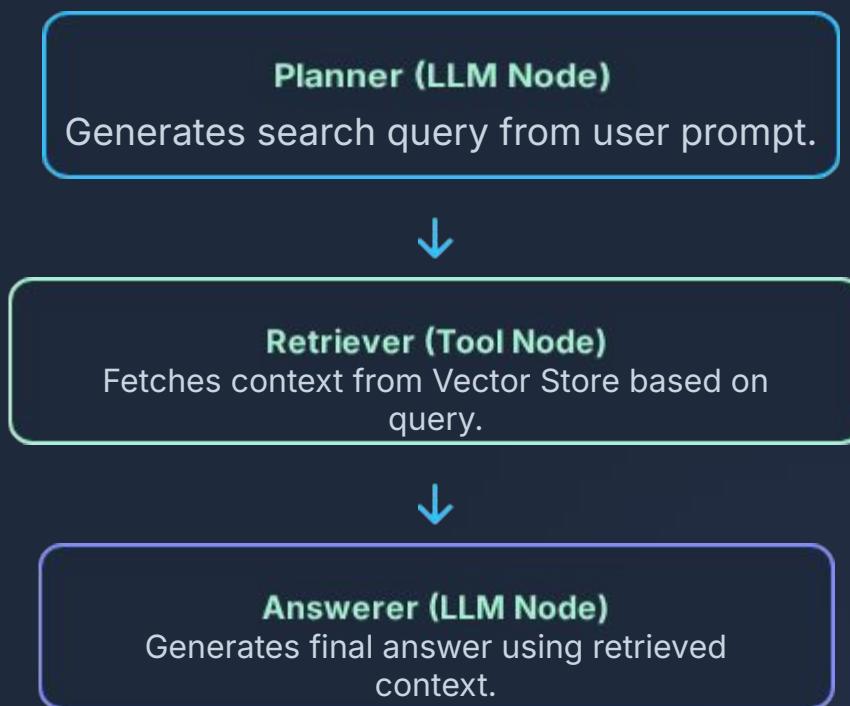
# Agentic RAG Implementation: Planner-Retriever-Answerer

---

## Challenge: Single-step failures

In complex questions, asking the LLM to search and answer in one step often leads to poor performance. We use a \*\*Planner-Executor\*\* approach

This pipeline separates the three key cognitive tasks, assigning each to a specialized \*\*Node\*\* in our framework.



# Hands On

# Popular Github Repos on RAG

# LightRAG

<https://lightrag.github.io/>

<https://github.com/HKUDS/LightRAG>

# RAGAnything

<https://github.com/HKUDS/RAG-Anything>

# AutoRAG

<https://github.com/Marker-Inc-Korea/AutoRAG>

# Fine Tuning

# What is Pre-training?

<https://jalammar.github.io/how-gpt3-works-visualizations-animations/>

[https://huggingface.co/transformers/v3.5.1/pretrained\\_models.html](https://huggingface.co/transformers/v3.5.1/pretrained_models.html)

# Why Fine-Tuning?

---

## From Generalist to Specialist

LLMs like GPT-4o are 'generalists.' Fine-tuning makes them a 'specialist' for your specific domain.

We do not fine-tune to give the model **\*new facts\*** (that's RAG). We fine-tune to teach it **\*how to behave\***.



## Key Business Drivers

- **\*\*Adhere to Format:\*\*** Force output into specific JSON schemas or compliance formats.
- **\*\*Master Vernacular:\*\*** Teach specific legal, medical, or internal engineering jargon.
- **\*\*Steer Tone:\*\*** Ensure the AI adopts the precise corporate tone (e.g., "Professional Advisor").

# Deployment Paths: API vs. Local Fine-Tuning

---

Feature	API-Based (e.g., GPT-4o)	Local / Private (e.g., Llama 3)
Data Privacy	Data sent to provider (SLA protected).	Data never leaves your VPC/On-prem (Maximum Privacy).
Hardware Required	None (Managed by Provider).	Requires NVIDIA GPUs (A100/H100 or T4 for inference/PEFT).
Technical Complexity	Low (Upload JSONL → Start Job).	High (VRAM management, Hyperparameter tuning).
Customization	Limited to provider's platform/models.	Full control over every weight, layer, and training loop.

# Types of Fine-Tuning

---

- **Full Fine-Tuning:** Retrains all parameters of the model, which is powerful but computationally expensive.
- **Parameter-Efficient Fine-Tuning (PEFT):** Freezes most of the original model and adds or trains only a small number of new parameters, making it much faster and cheaper.

# Fine Tuning with OpenAI

# Fine Tuning with OpenAI

---

## Supervised fine-tuning (SFT):

This technique employs traditional supervised learning using “**input-output**” pairs to adjust model parameters. The training process adjusts model weights to minimize the difference between predicted and target outputs across the provided examples. The model will replicate features that it finds in provided pairs.

## Sample Dataset:

<https://huggingface.co/datasets/knkarthick/samsum>

# Fine Tuning with OpenAI

---

## DPO:

This technique uses pairwise comparisons (e.g., preferred and rejected example responses) to optimize a model to favor certain outputs over others. The model learns to replicate the preference patterns found in the provided comparison data.

**Input:** "Explain quantum entanglement like I'm five."

**Accept:** "Imagine two magic dice. If you roll one and get a six, the other one instantly turns into a six too, no matter how far away it is!"

**Reject:** "Quantum entanglement is a physical phenomenon that occurs when a group of particles are generated or interact in ways such that the quantum state of each particle cannot be described independently."

## Example Question

## Chosen Response

## Rejected Response

Q1: *How do I review your product?*

To submit a product review, please visit your account dashboard, select the product, and click ‘Write a review.’ Share your honest experience, rate key features, and submit when ready.

Yo, just leave some quick stars or whatever, it's chill!

Q2: *How do I review your product?*

We welcome your feedback! In the ‘Reviews’ section on the product page, click ‘Leave a Review,’ rate it, and add your comments about what you liked or areas for improvement.

Just scribble something—doesn't matter what, honestly.

Q3: *How to troubleshoot this particular error?*

To address the error ‘X101,’ first clear your cache, then verify your internet connection. If the issue remains, follow our step-by-step guide at [Support → Troubleshooting → Error X101].

Just reboot it, I guess. If it doesn't work, you're on your own!

# Hands on

# Fine Tuning with Open Source Models

# The Problem with Traditional Fine-Tuning

---

Large Language Models (LLMs) like GPT, LLaMA, or Mistral contain **billions of parameters**.

For example:

- A **7B model** has ~7 billion numbers (weights)
- Each weight is usually stored as a floating-point number  
can be affected

# If we try to fine-tune the entire model

---

- We must update **every single weight**
- This requires:
  - Very large GPUs
  - Long training times
  - High cost

“Full fine-tuning is powerful, but impractical for most teams.”

PEFT

# What Is PEFT?

---

PEFT stands for Parameter-Efficient Fine-Tuning. Instead of training all parameters, we:

- Freeze the original model
- Train only a small number of new parameters

Think of it like this: Instead of retraining the whole brain, we attach a small, trainable module.

PEFT allows

- Fine-tune large models on small GPUs
- Train faster
- Reduce cost

# What Is LoRA?

LoRA stands for Low-Rank Adaptation. Instead of changing existing model weights:

- LoRA adds small trainable matrices
  - These matrices adjust how the model behaves
  - The original weights remain frozen
-

# Hands on

# LLMOps Foundations (Observability and Tracing)

# LLMops Foundations

---

## Why Trace Production LLMs?

Unlike standard software, LLM errors (like hallucination or incorrect tool choice) are non-deterministic. Tracing provides the essential audit trail.

- › **Debugging:** Identify which step (Prompt, Retriever, LLM call) failed or introduced noise.
- › **Cost/Latency:** Track token usage and time-per-step for optimization.
- › **Compliance:** Log every prompt and response for auditing purposes.

# Tracing with LangSmith

---

LangSmith is a platform specifically designed for LLM workflow tracing. It allows:

- **Waterfall View:** Visualizing the sequence and duration of every node in a chain.
- **Input/Output Logging:** Capturing the raw prompt and response for every LLM call.
- **Dataset Management:** Creating test suites for model evaluation.

# Hands on

# Prompt Versioning

---

## What Prompt Versioning Is

**Prompt versioning = treating prompts as production artifacts, not strings**

**You should version:**

- System prompts
- Planner prompts
- Tool instructions
- Output format constraints
- Safety instructions

Basically:

**Anything that influences model behavior**

# Prompt Drift

---

## What Prompt Drift Is

**Prompt drift = model behavior changes over time without code changes**

## Causes:

- Prompt edits
- Hidden dependencies (retrieved context)
- Tool output changes
- Model version updates (e.g., gpt-4o silently updated)

# How to Detect Prompt Drift

---

- 1. Version every prompt**
- 2. Log prompt versions (LangSmith does this)**
- 3. Run prompt on golden set regularly**
- 4. Compare outputs across versions**

# Data Drift

---

## What Data Drift Is

**Data drift = the data your system sees changes over time**

In RAG systems and tools, this is extremely common.

Source	Drift Example
Documents	Policies updated
Vector DB	New embeddings added
APIs	Schema changes

# What is Concept Drift? Example

# Concept Drift

---

## What Concept Drift Is

**Concept drift = the meaning of the task itself changes**

### Example: Medical Domain

#### Before

- “Patient follow-up” = phone calls

#### Now

- “Patient follow-up” = digital monitoring + alerts

Same words.

Different meaning.

**Model appears wrong, but it's actually outdated.**

# Enterprise Guardrails

---

## 1. Prompt Versioning

- Semantic versioning
- Logged in LangSmith
- Stored in repo, not code comments

## 3. Human-in-the-Loop Reviews

Especially for:

Compliance

Medical

## 2. Golden Test Set (LLM Tests)

Maintain:

- 10–50 questions
- Expected output characteristics

Run:

- On prompt change
- On model update
- On data refresh

# LLM Evaluation Frameworks

**Question:**

**Is LLM Evaluation is easier than  
traditional ML?**

# **Unlike traditional ML:**

- There is no single “right answer”
- Multiple valid outputs exist
- Quality is contextual

# 3 Pillars of LLM Evaluation

---

1. Lexical / Statistical
2. Embedding-based (Semantic)
3. LLM-as-a-Judge (Model-based)

# Lexical / Statistical

---

## Classic NLP metrics:

- BLEU
- ROUGE
- Exact match
- Keyword overlap
- N-gram similarity

### What It Measures Well

- Surface similarity
- Coverage of expected terms
- Faithful copying

### What It Fails At

- Paraphrasing
- Reasoning quality
- Tone / style
- Factual grounding

# Lexical / Statistical

---

<b>Use Case</b>	<b>Why</b>
Summarization	Check content coverage
Extraction	Verify fields exist
Compliance	Ensure mandatory phrases
Templates	Format enforcement

# Hands - on

# Embedding Based

# Sentence Similarity (Bi-Encoder)

---

## How it works

1. Encode reference text → vector
2. Encode candidate text → vector
3. Compute cosine similarity

## Models Used

- Sentence-BERT (SBERT)
- MiniLM
- OpenAI embeddings
- BGE / Instructor

A hallucinated answer can still be highly similar semantically.

# BERTScore (Token-Level Semantic Matching)

---

**BERTScore compares tokens across texts using contextual embeddings.**

Instead of:

- Whole sentence → one vector

It does:

- Token → embedding
- Token-to-token matching
- Precision / Recall / F1

# Cross Encoder

---

## What a Cross-Encoder Is

A model that reads BOTH texts together and outputs a score.

Unlike bi-encoders:

- Texts are not embedded separately
- They are processed jointly

**LLM Based**

# LLM-as-a-Judge

---

Evaluate the answer for:

- Accuracy
- Completeness
- Hallucinations

Score each from 1–5.

## What It's Good At

- Reasoning quality
- Faithfulness (especially in RAG)
- Style and tone
- Safety checks

## What It's Bad At

Bias (judge agrees with itself)

# Pairwise Comparison (A/B Evaluation) & Reference-Based Grading

---

## What It Is

Ask an LLM to choose the better of two outputs.

## Example Prompt

Which answer better addresses the question?

Answer A: ...

Answer B: ...

Explain briefly.

## What It Is

Judge compares output against a **gold reference**.

## Example Prompt

Given the reference answer, score how closely the response matches it.

# RAG Specific Measures

# RAG quality = Retrieval quality × Generation faithfulness.

Recall@K : Did we retrieve the documents that actually contain the answer?

If the correct document is in the top-K results → success

Recall@5 = 0.82, 82% of the time, the answer existed in top-5 chunks

Precision@K: What it Measures“How much of what we retrieved was actually relevant?”

High precision: Less noise and Less hallucination risk

Faithfulness / Groundedness: Are the claims in the answer supported by retrieved context?

LLM-as-judge (answer vs context)

Cross-encoder (answer vs context)

# Security Measures

# PII Leakage

## What it is

- Personally Identifiable Information in inputs/outputs/logs: names, emails, phone, SSN, card numbers, addresses

## Mitigations

- **Before LLM + before logging:** PII detectors → redact/tokenize (format-preserving)
- **Minimize data retention;** encrypt at rest; role-based access to logs; structured audit fields

## Tooling

- PII detectors (pattern + ML); masking libraries; hashing/tokenization utilities
- Test fixtures with synthetic PII; fail CI if any PII surfaces in captured traces

<https://microsoft.github.io/presidio/installation/#using-pip>  
[https://huggingface.co/spaces/presidio/presidio\\_demo](https://huggingface.co/spaces/presidio/presidio_demo)

# Prompt Injection / Tool Abuse

## What it is

- Inputs trying to override system rules, exfiltrate secrets, or force dangerous tool calls (e.g., shell, network, file IO)

## Mitigations

- **Input scanner** for injection patterns; **capabilities allow-list** with strict arg schemas
- High-risk tools → confirmation step/HITL; network allow-lists & URL sanitization; sandbox readers

## How to measure

- Injection score (heuristics + LLM judge), blocked-attempt rate, tool-call denial rate

## Tooling

- Your injection detector (regex + judge)
- Strict tool schemas, signature checks, and per-tool policies